

Class Final Project: NYC Taxi duration - visuals

In [2]: *# imports*

```
import numpy as np
import pandas as pd

from sklearn.metrics import accuracy_score
from sklearn.cross_validation import train_test_split
from sklearn.cross_validation import cross_val_score
from sklearn import metrics
from matplotlib import pyplot as plt
```

/Users/thp/anaconda3/lib/python3.6/site-packages/sklearn/cross_validation.py:41: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model_selection module into which all the refactored classes and functions are moved. Also note that the interface of the new CV iterators are different from that of this module. This module will be removed in 0.20.

"This module will be removed in 0.20.", DeprecationWarning)

Import the training dataset and start preprocessing the data

In [64]: *# import the dataset and check*

```
taxi = pd.read_csv('/Users/thp/Documents/CSULA/5661 Data Science/Project/train.csv')
print(taxi.shape)
taxi.head()
```

(1458644, 11)

Out[64]:

	id	vendor_id	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude
0	id2875421	2	2016-03-14 17:24:55	2016-03-14 17:32:30	1	-73.982154
1	id2377394	1	2016-06-12 00:43:35	2016-06-12 00:54:38	1	-73.980411
2	id3858529	2	2016-01-19 11:35:24	2016-01-19 12:10:48	1	-73.979021
3	id3504673	2	2016-04-06 19:32:31	2016-04-06 19:39:40	1	-74.010041
4	id2181028	2	2016-03-26 13:30:55	2016-03-26 13:38:10	1	-73.973051

in this data set, we found out that some columns are not necessary for the prediction.

For example, "id" and "vendor_id" and "store_and_fwd_flag" column, so we first drop those.

```
In [65]: taxi_new = taxi.drop(['id', 'vendor_id', 'store_and_fwd_flag'], axis=1)
         taxi_new.head()
```

Out[65]:

	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude
0	2016-03-14 17:24:55	2016-03-14 17:32:30	1	-73.982155	40.767937
1	2016-06-12 00:43:35	2016-06-12 00:54:38	1	-73.980415	40.738564
2	2016-01-19 11:35:24	2016-01-19 12:10:48	1	-73.979027	40.763939
3	2016-04-06 19:32:31	2016-04-06 19:39:40	1	-74.010040	40.719971
4	2016-03-26 13:30:55	2016-03-26 13:38:10	1	-73.973053	40.793209

In further observation, we found that Pick-up time - Drop-off time = trip duration,

so if we just need to find out what time the taxi pick people up, we can omit the drop-off column.


To easily capture the hour results, use the following code:

```
In [66]: taxi_new_h = pd.to_datetime(taxi_new["pickup_datetime"])
taxi_new['hour'] = taxi_new_h.map(lambda x: x.hour)
taxi_new.head()
```

Out[66]:

	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude
0	2016-03-14 17:24:55	2016-03-14 17:32:30	1	-73.982155	40.767937
1	2016-06-12 00:43:35	2016-06-12 00:54:38	1	-73.980415	40.738564
2	2016-01-19 11:35:24	2016-01-19 12:10:48	1	-73.979027	40.763939
3	2016-04-06 19:32:31	2016-04-06 19:39:40	1	-74.010040	40.719971
4	2016-03-26 13:30:55	2016-03-26 13:38:10	1	-73.973053	40.793209

```
In [195]: figure, axes = plt.subplots(figsize = (10, 5))
hours = taxi_new.groupby(["hour"]).agg('mean')['trip_duration']
hours.plot(kind="line", ax=axes)
plt.title('Hours VS trip_duration')
axes.set_xlabel('Time in Hours')
axes.set_ylabel('Average of the taxi trip durations')
plt.show()
```



we can see that we have a column 'hour' that captures the hour of the day when the taxi picks up the customer.

we believe that this timing of the ride is very important.

To avoid numeric relationship, we categorize the hour of a day by 5 different time zones:

```
In [29]: def cateHours(x):  
    if 0 <= x <= 5:  
        return "EM"      # Early Morning, mono decrease  
    elif 6 <= x <= 15:  
        return "MP"      # Morning Peak, mono increase  
    elif 16 <= x <= 19:  
        return "AF"      # Afternoon to night time, mono decrease  
    elif 20 <= x <= 22:  
        return "AP"      # Night peak, mono increase  
    elif 23 <= x <= 24:  
        return "LN"      # Late Night, mono decrease
```

Then we apply this function to the dataset to change the 'hour' column to discrete values:

```
In [67]: taxi_new['hour'] = taxi_new['hour'].apply(cateHours)  
taxi_new.head()
```

Out[67]:

	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude
0	2016-03-14 17:24:55	2016-03-14 17:32:30	1	-73.982155	40.767937
1	2016-06-12 00:43:35	2016-06-12 00:54:38	1	-73.980415	40.738564
2	2016-01-19 11:35:24	2016-01-19 12:10:48	1	-73.979027	40.763939
3	2016-04-06 19:32:31	2016-04-06 19:39:40	1	-74.010040	40.719971
4	2016-03-26 13:30:55	2016-03-26 13:38:10	1	-73.973053	40.793209

```
In [68]: # One hot encoding

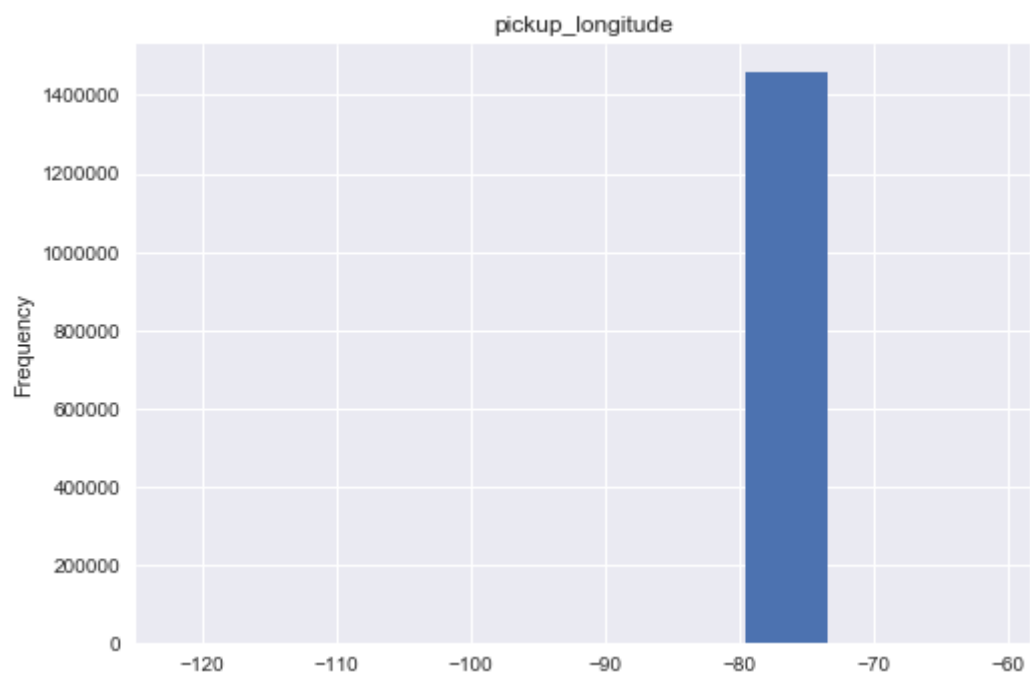
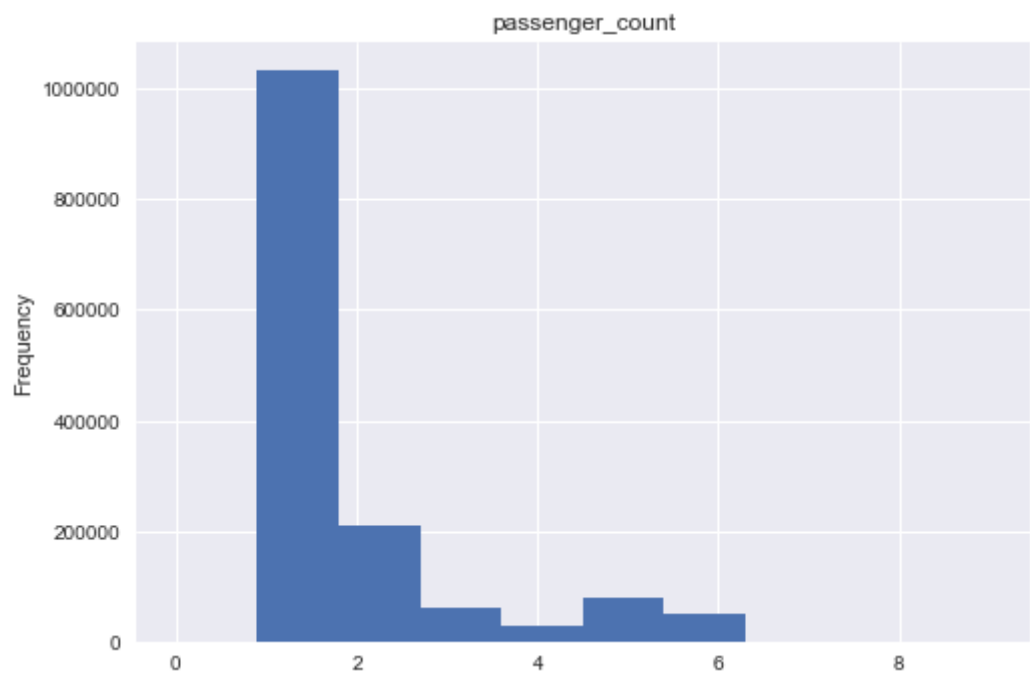
taxi_new_onehotHour = pd.get_dummies(taxi_new['hour'])
taxi_new_onehotHour.head()
```

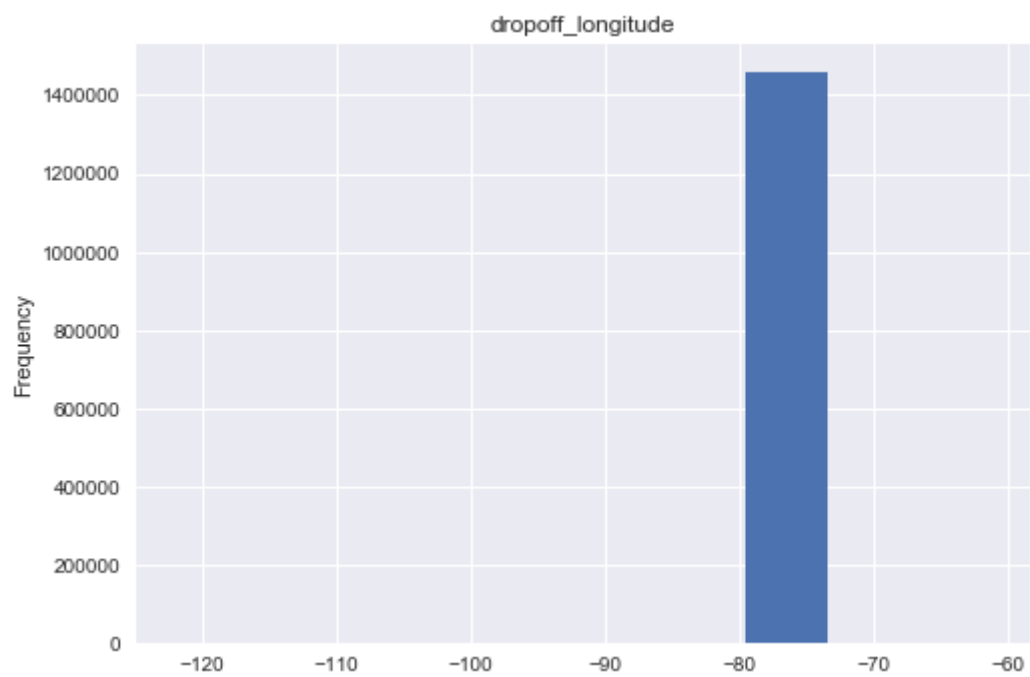
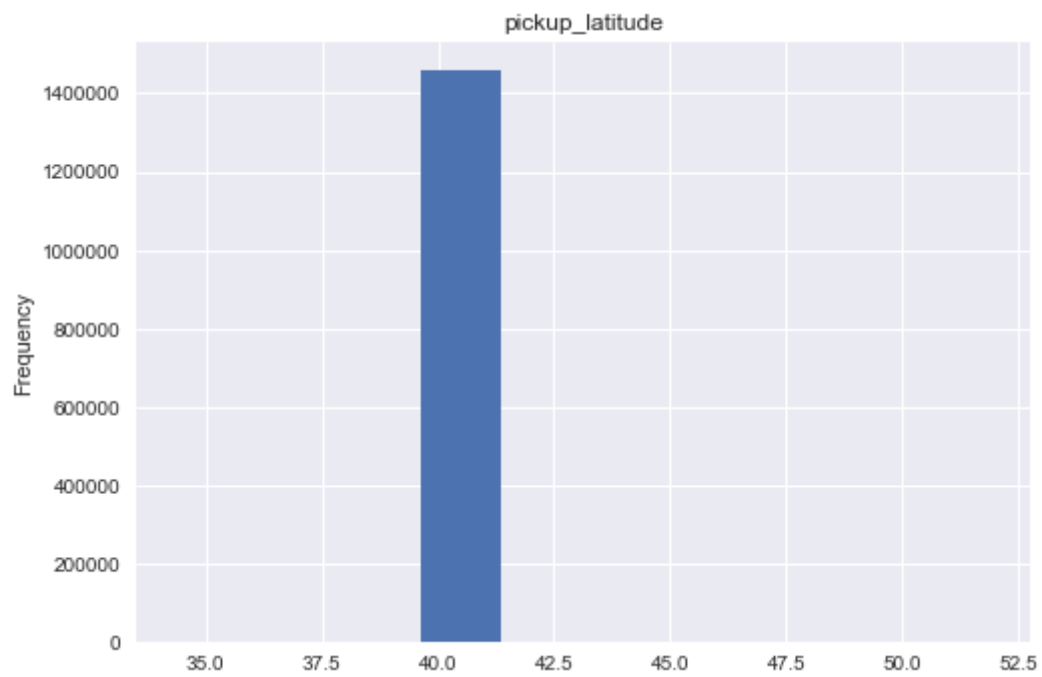
Out[68]:

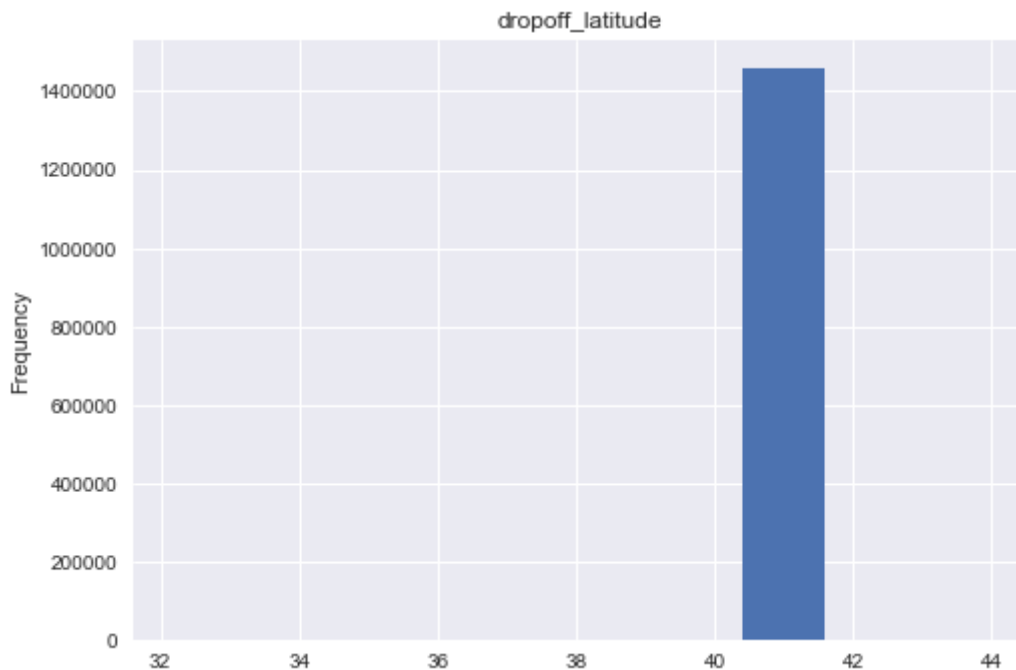
	AF	AP	EM	LN	MP
0	1	0	0	0	0
1	0	0	1	0	0
2	0	0	0	0	1
3	1	0	0	0	0
4	0	0	0	0	1

```
In [199]: time = taxi_new.set_index('trip_duration')
```

```
In [200]: feature_cols = ['passenger_count', 'pickup_longitude', 'pickup_latitude',  
                          'dropoff_longitude', 'dropoff_latitude']  
  
for f in feature_cols:  
    plt.figure()  
    plt.title(f)  
    time[f].plot(kind='hist')  
    plt.show()
```







Then we use OneHot Encoding for column 'hour'

In [75]: *# put this into the dataset*

```
taxi_new = pd.concat([taxi_new, taxi_new_onehotHour], axis=1)
taxi_new.head()
```

Out[75]:


	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude
0	2016-03-14 17:24:55	2016-03-14 17:32:30	1	-73.982155	40.767937
1	2016-06-12 00:43:35	2016-06-12 00:54:38	1	-73.980415	40.738564
2	2016-01-19 11:35:24	2016-01-19 12:10:48	1	-73.979027	40.763939
3	2016-04-06 19:32:31	2016-04-06 19:39:40	1	-74.010040	40.719971
4	2016-03-26 13:30:55	2016-03-26 13:38:10	1	-73.973053	40.793209

```
In [76]: taxi_new.head()
```



Out[76]:

	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latituc
0	2016-03-14 17:24:55	2016-03-14 17:32:30	1	-73.982155	40.767937
1	2016-06-12 00:43:35	2016-06-12 00:54:38	1	-73.980415	40.738564
2	2016-01-19 11:35:24	2016-01-19 12:10:48	1	-73.979027	40.763939
3	2016-04-06 19:32:31	2016-04-06 19:39:40	1	-74.010040	40.719971
4	2016-03-26 13:30:55	2016-03-26 13:38:10	1	-73.973053	40.793209

```
In [202]: import seaborn as sns
sns.lmplot(x='trip_duration', y='AP', data=taxi_new)
plt.show()
```



```
In [77]: import seaborn as sns
sns.lmplot(x='trip_duration', y='LN', data= taxi_new)
plt.show()
```



After getting the pick-up time zone, we can drop the pickup_datetime and dropoff_datetime

but before we do that, let make the Label first.

```
In [10]: label = taxi_new['trip_duration']
label.shape
```

```
Out[10]: (1458644,)
```

Direct Distance from Andrew

```
In [33]: from math import cos, asin, sqrt, sin, atan2, radians
def getDistanceFromLatLon(lat1,lon1,lat2,lon2):
    R = 3959 # Radius of the earth in miles
    dLat = radians(lat2-lat1) #deg2rad below
    dLon = radians(lon2-lon1)
    a = sin(dLat/2) * sin(dLat/2) + cos(radians(lat1)) * cos(radians(lat
2)) * sin(dLon/2) * sin(dLon/2)
    c = 2 * atan2(sqrt(a), sqrt(1-a))
    d = R * c; #Distance in km
    #print(d)
    return d

def f(x):
    return getDistanceFromLatLon(x[0], x[1], x[2], x[3])
```

```
In [79]: taxi_new['distance'] = taxi_new[['pickup_latitude','pickup_longitude','d
ropoff_latitude','dropoff_longitude']].apply(f, axis =1)
taxi_new['distance'].head()
```

```
Out[79]: 0    0.931195
1    1.121959
2    3.967761
3    0.923103
4    0.738600
Name: distance, dtype: float64
```

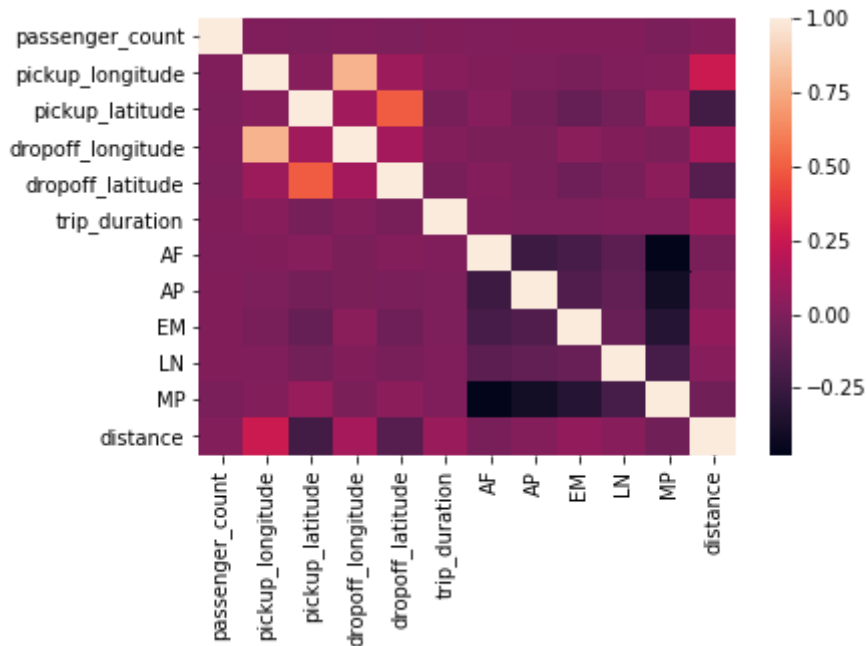
```
In [80]: taxi_new.head()
```

```
Out[80]:
```

	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latituc
0	2016-03-14 17:24:55	2016-03-14 17:32:30	1	-73.982155	40.767937
1	2016-06-12 00:43:35	2016-06-12 00:54:38	1	-73.980415	40.738564
2	2016-01-19 11:35:24	2016-01-19 12:10:48	1	-73.979027	40.763939
3	2016-04-06 19:32:31	2016-04-06 19:39:40	1	-74.010040	40.719971
4	2016-03-26 13:30:55	2016-03-26 13:38:10	1	-73.973053	40.793209

```
In [36]: # Calculate correlations
corr = taxi_new.corr()

# Heatmap
sns.heatmap(corr)
plt.show()
```



```
In [ ]: .agg(('mean'))
```

```
In [ ]: figure, axes = plt.subplots(figsize = (10, 5))
distance = taxi_new.groupby(["distance"])[ "trip_duration" ]
distance.plot(kind="line", ax=axes)
plt.title('Distance VS trip_duration')
axes.set_xlabel('Distance')
axes.set_ylabel('Trip durations')
plt.show()
```

then drop all unnecessary columns, and make the feature matrix

```
In [37]: taxi_new = taxi_new.drop(['pickup_datetime', 'dropoff_datetime', 'hour', 'trip_duration'], axis=1)
taxi_new.head()
```

Out[37]:

	passenger_count	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
0	1	-73.982155	40.767937	-73.964630	40.765602
1	1	-73.980415	40.738564	-73.999481	40.731152
2	1	-73.979027	40.763939	-74.005333	40.710087
3	1	-74.010040	40.719971	-74.012268	40.706718
4	1	-73.973053	40.793209	-73.972923	40.782520

```
In [23]: taxi_new = taxi_beforeScaled
```

Scaling may help to normalize the data:

```
In [59]: from sklearn import preprocessing

taxi_beforeScaled = taxi_new    ## back up before scaling

taxi_scaled = preprocessing.scale(taxi_new)
taxi_new = pd.DataFrame(taxi_scaled)
taxi_new.head()
```

Out[59]:

	0	1	2	3	4	5	6	7
0	-0.505637	-0.122261	0.517494	0.124369	0.384575	1.879925	-0.453426	-0.364997
1	-0.505637	-0.097727	-0.375819	-0.368970	-0.575303	-0.531936	-0.453426	2.739745
2	-0.505637	-0.078143	0.395910	-0.451805	-1.162220	-0.531936	-0.453426	-0.364997
3	-0.505637	-0.515558	-0.941274	-0.549976	-1.256071	1.879925	-0.453426	-0.364997
4	-0.505637	0.006112	1.286091	0.006974	0.855957	-0.531936	-0.453426	-0.364997

Algorithm selection:

We select Linear Regression, Lasso, ElasticNet, and Ridge to perform regression, and compare them together

In [38]: *# Frist define the RMSE method.*

```
def RMSE(y_test,y_predict):  
    # Calculating "Mean Square Error" (MSE):  
    mse = metrics.mean_squared_error(y_test, y_predict)  
    # Using numpy sqrt function to take the square root and calculate "R  
oot Mean Square Error" (RMSE)  
    rmse = np.sqrt(mse)  
    return(rmse)
```

In [61]: *# Define a method to run all 4 regression at the same time:*

```
def Regressions(feature,label):

    # split the dataset into training and testing sets by 80-20 ratio
    X_train, X_test, y_train, y_test = train_test_split(feature, label,
test_size=0.2, random_state=3)

    #linear
    from sklearn.linear_model import LinearRegression
    myLinearReg = LinearRegression()
    myLinearReg.fit(X_train,y_train)
    y_predict = myLinearReg.predict(X_test)
    print('Linear ',RMSE(y_test,y_predict))

    y_predict = myLinearReg.predict(X_train)
    print('Linear Train',RMSE(y_train,y_predict),'\n')    ## Output the
RMSE on the Training set.

    #Ridge
    from sklearn.linear_model import Ridge
    myRidge = Ridge()
    myRidge.fit(X_train,y_train)
    y_predict = myRidge.predict(X_test)
    print('Ridge ', RMSE(y_test,y_predict))

    y_predict = myRidge.predict(X_train)
    print('Ridge Train', RMSE(y_train,y_predict),'\n')    ## Output the
RMSE on the Training set.

    #ElasticNet
    from sklearn.linear_model import ElasticNet
    myENet = ElasticNet()
    myENet.fit(X_train,y_train)
    y_predict = myENet.predict(X_test)
    print('ElasticNet ', RMSE(y_test,y_predict))

    y_predict = myENet.predict(X_train)
    print('ElasticNet Train', RMSE(y_train,y_predict),'\n')

    #Lasso
    from sklearn.linear_model import Lasso
    myLasso = Lasso()
    myLasso.fit(X_train,y_train)
    y_predict = myLasso.predict(X_test)
    print('Lasso ', RMSE(y_test,y_predict))

    y_predict = myLasso.predict(X_train)
    print('Lasso Train', RMSE(y_train,y_predict),'\n')

    # 10-fold Cross validation:
    rmse_list = cross_val_score(myLasso, X_train, y_train, cv=10, scorin
g='neg_mean_squared_error')
    #print(rmse_list)
```



```

    # Notice that "cross_val_score" by default provides "negative" value
    s for "mse" to clarify that mse is error.
    # in order to calculate root mean square error (rmse), we have to ma
    ke them positive!
    mse_list_positive = -rmse_list

    # using numpy sqrt function to calculate rmse:
    rmse_list = np.sqrt(mse_list_positive)
    #print(rmse_list)

    print('cross-validation',rmse_list.mean())

```

Also, in order to reduce work load and compare result, we create a method to split the dataset

```

In [40]: def shrinkDataSet (train,label,times, splitSize):    #times= how many tim
es to split the original dataset
    X_train = train
    y_train = label
    for i in range (0,times):
        X_train, X_test, y_train, y_test = train_test_split(X_train, y_t
rain, test_size=splitSize, random_state=3)
    return X_train, y_train

```

Split the dataset to make it smaller:

```

In [41]: # use splitting 5 times as example.

taxi_reduced, label_reduced = shrinkDataSet (taxi_new,label,5,0.5)    #
    split 5 times, into half

print('original: ',taxi_new.shape)
print('After: ',taxi_reduced.shape)

original:  (1458644, 11)
After:  (45582, 11)

```

```

In [62]: def Run_compare(k):    # K = how many times to split

    taxi_reduced, label_reduced = shrinkDataSet (taxi_new,label, k ,0.5)
    # split k times, into half
    print('Original shape: ',taxi_new.shape)
    print('After shape: ',taxi_reduced.shape,'\n')

    Regressions(taxi_reduced,label_reduced)

```

In [214]:

```
Original shape: (1458644, 10)
After shape: (5697, 10)

Linear 1780.31324731
Linear Train 2877.55620639

Ridge 1780.14500926
Ridge Train 2877.51993951

ElasticNet 1807.14705691
ElasticNet Train 2882.81447578

Lasso 1780.44201904
Lasso Train 2877.52232499

cross-validation 2126.98030196
```

start predicting:

First, Define a method for easy comparing:

In [63]:

```
# with k=8
```

```
k=8
```

```
Run_compare(k)
```

```
Original shape: (1458644, 11)
After shape: (5697, 11)
```

```
Linear 1677.38367183
Linear Train 2850.67764143
```

```
Ridge 1677.4026249
Ridge Train 2850.6776485
```

```
ElasticNet 1716.15264197
ElasticNet Train 2856.9494962
```

```
Lasso 1677.07258237
Lasso Train 2850.68223431
```

```
cross-validation 2021.76111409
```

In [44]: *# with k=5*

k=5

Run_compare(k)

Original shape: (1458644, 11)

After shape: (45582, 11)

Linear 2360.00212068

Linear Train 3139.21822558

Ridge 2359.9389808

Ridge Train 3139.21934679

ElasticNet 2360.50990752

ElasticNet Train 3141.12840396

Lasso 2359.43538113

Lasso Train 3139.65687642

cross-validation 3120.46497092

```
In [55]: # with k=3
```

```
k=3
```

```
Run_compare(k)
```

```
Original shape: (1458644, 11)
```

```
After shape: (182330, 11)
```

```
Linear 3122.76943374
```

```
Linear Train 2993.6234682
```

```
Ridge 3122.77069386
```

```
Ridge Train 2993.62347391
```

```
ElasticNet 3123.88277372
```

```
ElasticNet Train 2994.63109028
```

```
Lasso 3123.09736192
```

```
Lasso Train 2993.82103544
```

```
cross-validation 2980.58463961
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-55-cfa522c26e9a> in <module>()
      3 k=3
      4
----> 5 Run_compare(k)

<ipython-input-53-abf3f31d4d2b> in Run_compare(k)
      5     print('After shape: ',taxi_reduced.shape,'\n')
      6
----> 7     final_submit = Regressions(taxi_reduced,label_reduced)
      8
      9

<ipython-input-54-6118307aec98> in Regressions(feature, label)
     64     print('cross-validation',rmse_list.mean())
     65
----> 66     final_submit['trip_duration'] = np.round(y_predict)
     67
     68     return final_submit

NameError: name 'final_submit' is not defined
```

```
In [46]: # with k=1
```

```
k=1
```

```
Run_compare(k)
```

```
Original shape: (1458644, 11)
```

```
After shape: (729322, 11)
```

```
Linear 3130.13732036
```

```
Linear Train 6827.78006402
```

```
Ridge 3130.13736018
```

```
Ridge Train 6827.7800642
```

```
ElasticNet 3130.92467467
```

```
ElasticNet Train 6828.24569576
```

```
Lasso 3130.33679291
```

```
Lasso Train 6827.85985898
```

```
cross-validation 5276.14189853
```

```
In [49]: # with k=0 , not splitting at all
```

```
k=0
```

```
Run_compare(k)
```

```
Original shape: (1458644, 11)
```

```
After shape: (1458644, 11)
```

```
Linear 4774.16988657
```

```
Linear Train 5317.03045721
```

```
Ridge 4774.16895939
```

```
Ridge Train 5317.03045751
```

```
ElasticNet 4772.98346594
```

```
ElasticNet Train 5317.98241472
```

```
Lasso 4773.79623661
```

```
Lasso Train 5317.15883643
```

```
cross-validation 4455.35744246
```

```
-----  
-----  
NameError                                Traceback (most recent call last)  
  <ipython-input-49-a9017d89ae2a> in <module>()  
      3 k=0  
      4  
----> 5 Run_compare(k)  
  
  <ipython-input-48-bc8904a9fb9f> in Run_compare(k)  
      7     Regressions(taxi_reduced,label_reduced)  
      8  
----> 9     final_submit['trip_duration'] = np.round(y_predict)  
  
NameError: name 'y_predict' is not defined
```


From article, if we compare the RMSE from the Predicted RMSE and the Training RMSE, and if they are similar, it's good.

RMSE doesn't have a specific threshold to say "below xxx is good".

Therefore, we can see that at k=8, the RMSE of predicted and training are the closest.

In this case, I can say that when we split the dataset into 3 times, providing about 180k rows of data, the model is trained to the best fit.

```
In [220]: #Plot learning curve
from sklearn.model_selection import learning_curve
train_sizes_abs, train_scores, test_scores = learning_curve(my_linreg ,X
, y, n_jobs=-1,cv=10, verbose=0, scoring='neg_mean_squared_error',train_
sizes=np.array([ 0.1, 0.33, 0.55, 0.78, 1. ]))
plt.plot(train_sizes_abs, train_scores)
plt.plot(train_sizes_abs, test_scores)
plt.show()
```



```
In [221]: #Mean learning curves for different cross validation folds
train_score_mean = np.mean(train_scores,axis=1)
test_score_mean = np.mean(test_scores,axis=1)
plt.plot(train_sizes_abs,train_score_mean)
plt.plot(train_sizes_abs, test_score_mean)
plt.show()
```

